

```
In [1]: from cartpole import *
import torch_utils as ptu
%matplotlib inline
plt.rcParams["figure.figsize"] = (10, 10)
```

```
In [2]: ptu.init_gpu(use_gpu=False) # use cpu

env = gym.make("CartPole-v1")
env.seed(42)

ptu.set_random_seed(42)

obs_dim = env.observation_space.shape[0]
act_dim = env.action_space.n

policy = ptu.MLPCategoricalpolicy(
    obs_dim, act_dim, hidden_sizes=[64, 64], activation="relu_inplace"
).to(ptu.device)
```

GPU not detected. Defaulting to CPU.

```
In [3]: lr = 3e-3
```

1


1. REINFORCE

Implement a vanilla reinforce algorithm given by the following gradient update for your policy.

```
In [4]: mean_reward_reinforce, train_loss_reinforce = reinforce(
    env,
    policy,
    num_itrs=200,
    batch_size=500,
    gamma=0.99,
    lr=lr)
```

100%|██████████| 200/200 [00:39<00:00, 5.09it/s]

```
In [5]: plot_loss_rew(train_loss_reinforce, mean_reward_reinforce)
```




1.2 Reward to go

Implement the policy gradient algorithm using update rule

```
In [6]: mean_reward_rtg, train_loss_rtg = vpg_with_baseline(
    env,
    policy,
    num_itr=200,
    batch_size=500,
    gamma=0.99,
    lr=lr,
    baseline=False,
)
plot_loss_rew(train_loss_rtg, mean_reward_rtg)
```

100%|██████████| 200/200 [00:41<00:00, 4.83it/s]




1.3 Reward to go + baseline

To reduce the variance of the estimated returns, subtract the returns using a constant b such that the mean of the modified returns is 0.

```
In [7]: mean_reward_baseline, train_loss_baseline = vpg_with_baseline(
    env,
    policy,
    num_itrs=200,
    batch_size=500,
    gamma=0.99,
    lr=lr,
    baseline=True,
)
plot_loss_rew(train_loss_baseline, mean_reward_baseline)
```

100%|██████████| 200/200 [00:39<00:00, 5.10it/s]



Although the final result collapse as well and it don't seems to be be than previous plots. The variance is smaller than reinforce and reward to go version. Each model should be tune with the associate hyperparameters. In order to fairly compare among 3 methods, I use same network size and depth and same learning rate.

With an good learning rate for baseline case, the result should be the best among 3 methodes.

```
In [8]: f, (ax1, ax2) = plt.subplots(2, 1)
ax1.plot(train_loss_reinforce, label= "Reinforce")
ax1.plot(train_loss_rtg, label="Reward to go")
ax1.plot(train_loss_baseline, label="Rtg + Baseline")

ax1.set_xlabel('Iteration')
ax1.set_ylabel('Loss')
ax1.set_title('Training loss at each iteration')
ax1.grid(True)
ax1.legend()

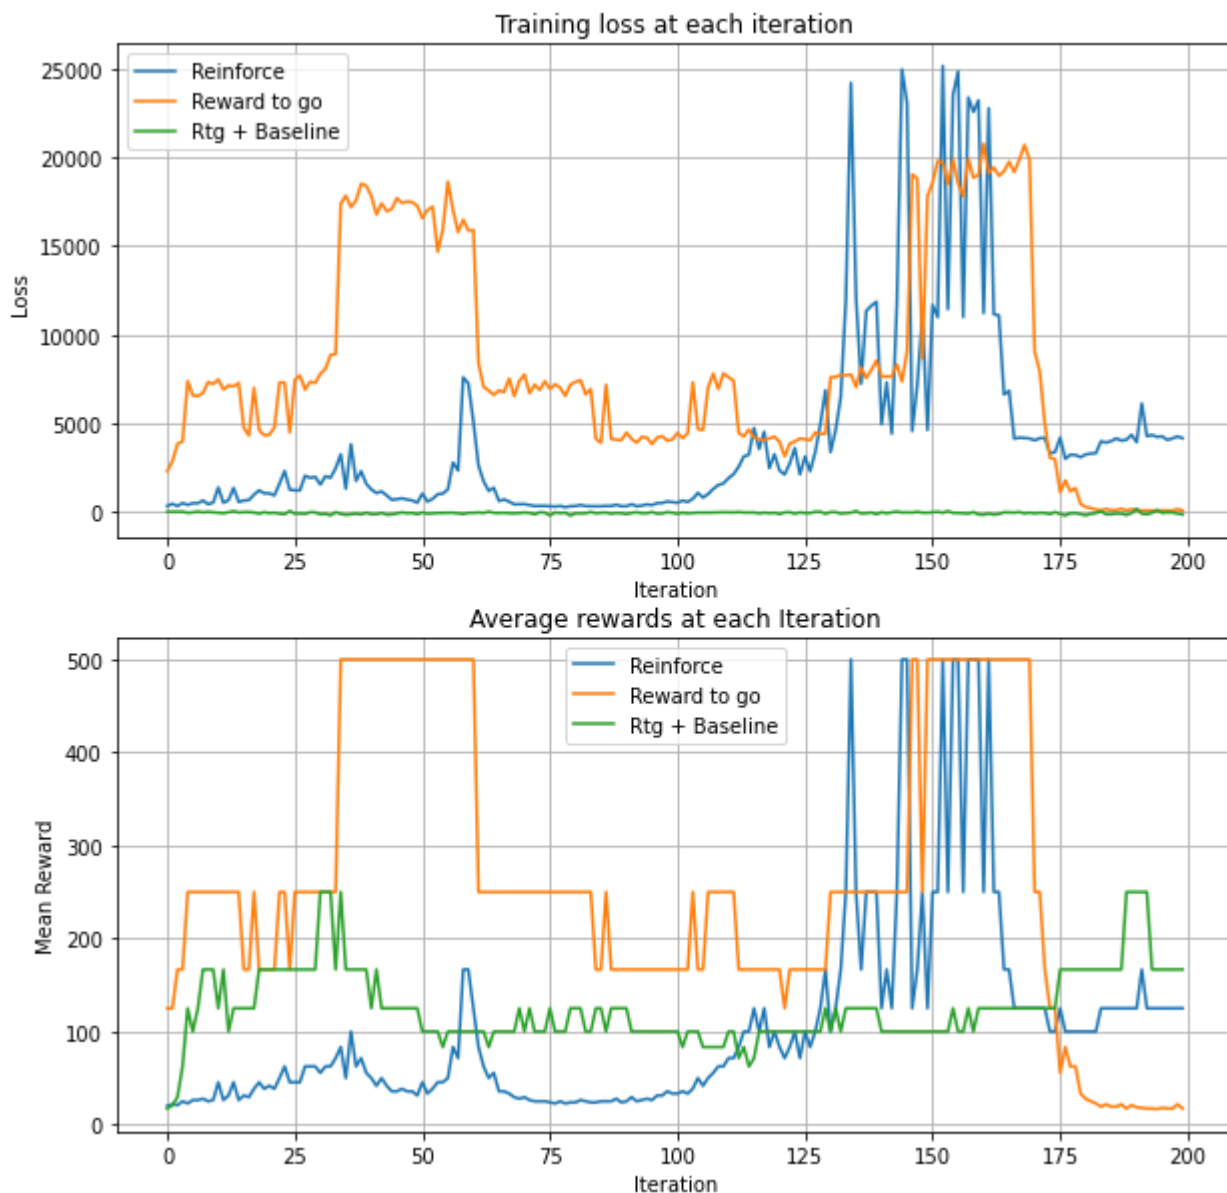
ax2.plot(mean_reward_reinforce, label= "Reinforce")
```

```

ax2.plot(mean_reward_rtg, label="Reward to go")
ax2.plot(mean_reward_baseline, label="Rtg + Baseline")

ax2.set_xlabel('Iteration')
ax2.set_ylabel('Mean Reward')
ax2.set_title('Average rewards at each Iteration')
ax2.grid(True)
ax2.legend()
plt.show()

```



In [16]:

```

# =====
# 600
lr = 7e-3
mean_reward_baseline_600, train_loss_baseline_600 = vpg_with_baseline(
    env,
    policy,
    num_itrs=200,
    batch_size=600,
    gamma=0.99,
    lr=lr,

```

```

        baseline=False,
    )

    mean_reward_baseline_800, train_loss_baseline_800 = vpg_with_baseline(
        env,
        policy,
        num_itrs=200,
        batch_size=800,
        gamma=0.99,
        lr=lr,
        baseline=False,
    )

    mean_reward_baseline_1000, train_loss_baseline_1000 = vpg_with_baseline(
        env,
        policy,
        num_itrs=200,
        batch_size=1000,
        gamma=0.99,
        lr=lr,
        baseline=False,
    )

```

```

100%|██████████| 200/200 [00:50<00:00, 3.99it/s]
100%|██████████| 200/200 [01:05<00:00, 3.05it/s]
100%|██████████| 200/200 [01:23<00:00, 2.39it/s]

```

In [13]:

```


f, (ax1, ax2) = plt.subplots(2, 1)
ax1.plot(train_loss_baseline_600, label="batch size: 600")
ax1.plot(train_loss_baseline_800, label="batch size: 800")
ax1.plot(train_loss_baseline_1000, label="batch size: 1000")

ax1.set_xlabel('Iteration')
ax1.set_ylabel('Loss')
ax1.set_title('Training loss at each iteration')
ax1.grid(True)
ax1.legend()

ax2.plot(mean_reward_baseline_600, label="batch size: 600")
ax2.plot(mean_reward_baseline_800, label="batch size: 800")
ax2.plot(mean_reward_baseline_1000, label="batch size: 1000")

ax2.set_xlabel('Iteration')
ax2.set_ylabel('Mean Reward')
ax2.set_title('Average rewards at each Iteration')
ax2.grid(True)
ax2.legend()
plt.show()

```



Since we are using MC estimation, increasing batch size did improve the training performnce.